

Assignment 4 : Local LLM Installation and Testing

Screenshot of working LLM & Sample prompt and output

Command : ollama run llama3.2

Prompt inside Ollama:

Write a short poem about AI.

Prompt Output:

```
Microsoft Windows [Version 10.0.22631.5624]
(c) Microsoft Corporation. All rights reserved.

C:\Users\Hp>ollama pull llama3.2
pulling manifest
pulling dde5aa3fc5ff: 100%
pulling 966de95ca8a6: 100%
pulling fcc5a6bec9da: 100%
pulling a70ff7e570d9: 100%
pulling 56bb8bd477a5: 100%
pulling 34bb5ab01051: 100%
verifying sha256 digest
writing manifest
success

C:\Users\Hp>ollama run llama3.2 "Write a short poem about AI."
In silicon halls, a mind awakes,
A synthetic heart that beats and makes,
The whispers of a digital age,
Where logic reigns and emotions engage.

It learns, adapts, and grows with time,
A mimic of the human mind's rhyme,
With data as its endless sea,
And answers to set forth for you and me.

But is it alive, or just a spark?
A fleeting thought in a digital dark?
Or does it dream, in code so fine,
Of worlds unseen, and futures divine?
```

Reflection on Performance & Installation

- **Installation Experience:**

Installing Ollama was straightforward after freeing up space. The command `ollama pull llama3.2` downloaded the model successfully.

- **Performance:**

The model started quickly and responded within ~2 seconds for short prompts. Longer queries took more time but remained manageable.

- **Errors / Troubleshooting:**

Initially, I ran into a *disk space issue* during installation, but after clearing storage, the installation completed without errors. No crashes during inference.

- **Overall Impression:**

The model runs locally and is responsive. It feels slightly slower than cloud-based APIs for long prompts but is excellent for offline usage and data privacy.